

Machine Learning Operations

Assignment 1



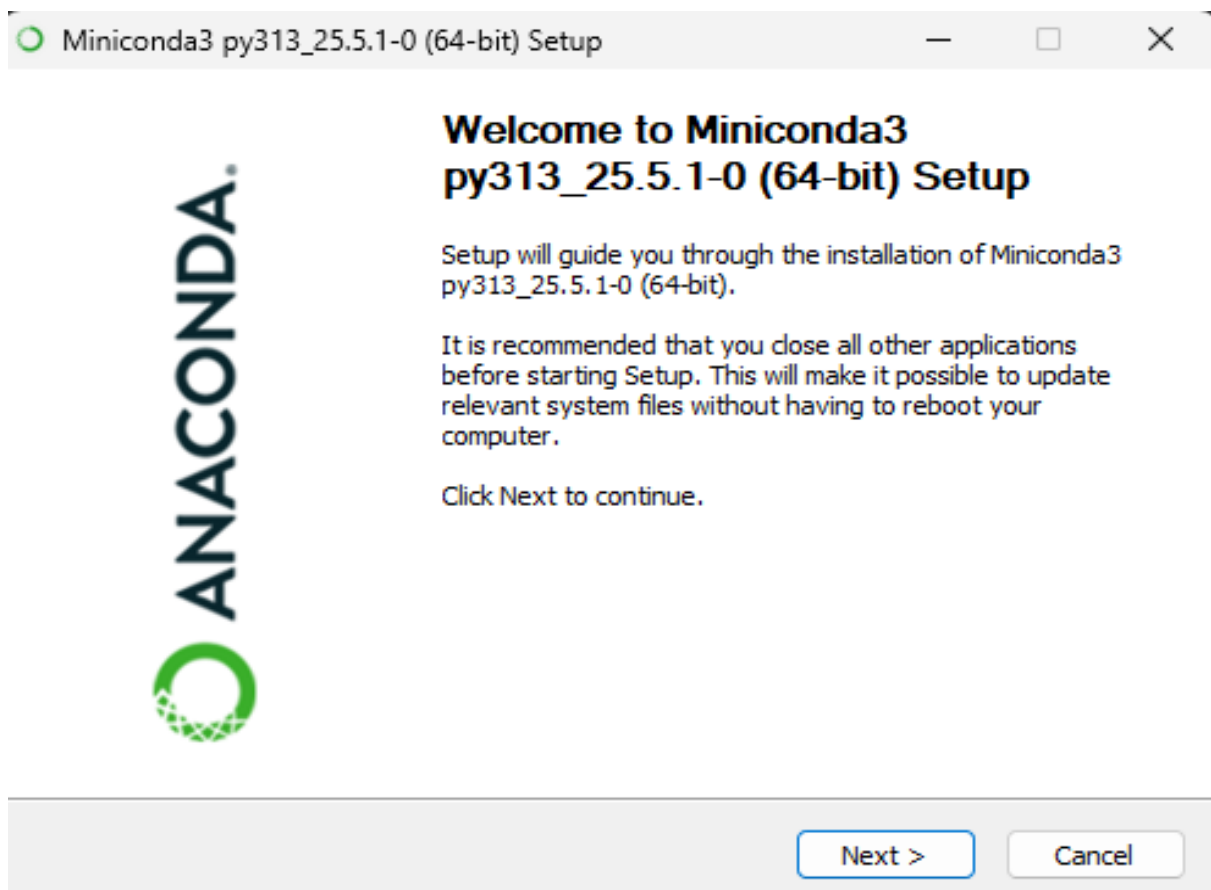
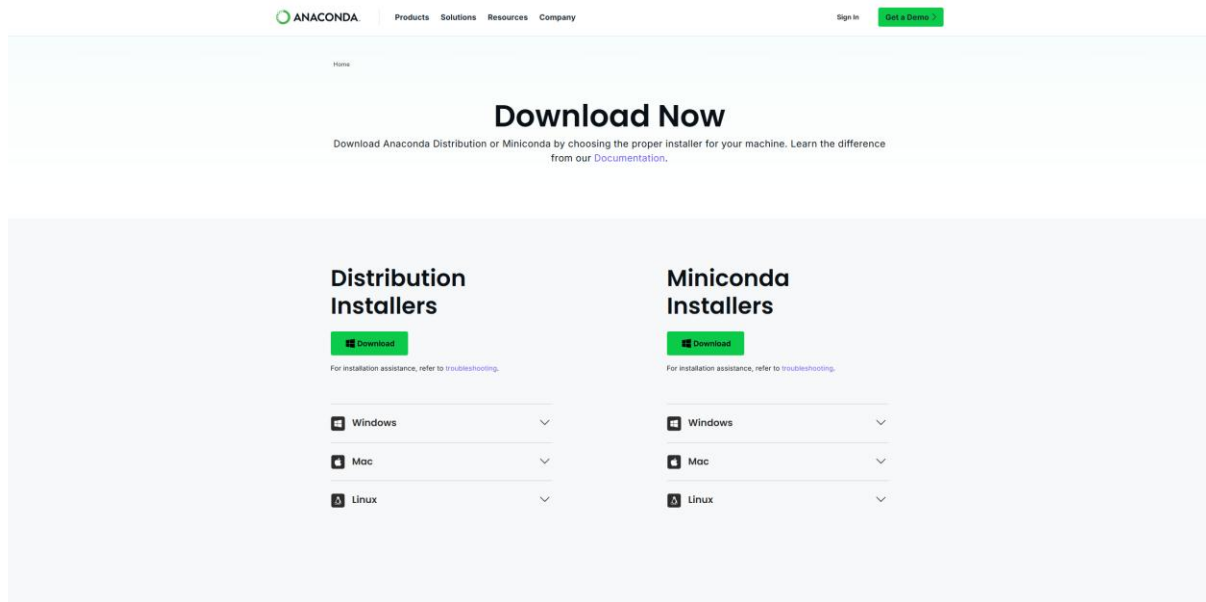
Submitted By – Akshay Kumar (G24AI1033)

GitHub Repository Link – [Github A1](#)

Conda setup

First step is to download Conda from the following link - [Download Now | Anaconda](#)

We follow the installation instructions provided on the website. We will run the downloaded installer and accept the default options.



Creating a Conda Environment for the Assignment

We will create a dedicated conda environment for the project. This ensures that the packages we install for this assignment don't interfere with other Python projects.

1. Create the environment:

conda create -n mlops_assignment1 python=3.9

```
(base) C:\Users\ASHAY>conda create -n mlops_assignment1 python=3.9
Channels:
 - defaults
Platform: win-64
Collecting package metadata (repodata.json): done
Solving environment: done

# Package Plan #

  environment location: C:\Users\ASHAY\miniconda3\envs\mlops_assignment1
  added / updated specs:
    - python=3.9

The following packages will be downloaded:



| package           | build         |         |
|-------------------|---------------|---------|
| python-3.9.21     | h716158d_0    | 16.7 MB |
| setuptools-70.1.1 | py39ha95532_0 | 1.7 MB  |
| wheel-0.43.1      | py39ha95532_0 | 145 KB  |
| Total:            |               | 18.5 MB |



The following NEW packages will be INSTALLED:



| package         | build                                 |
|-----------------|---------------------------------------|
| ca-certificates | ca-certificates-2025.2.25-haa95532_0  |
| certifi         | certifi-2025.2.25-haa95532_0          |
| libffi          | libffi-3.4.4-h8f7b12b_1               |
| openssl         | openssl-3.0.16-h8f7b12b_0             |
| pip             | pip-25.1.0-pyha772135_2               |
| python          | python-3.9.21-h716158d_0              |
| setuptools      | setuptools-70.1.1-py39ha95532_0       |
| sqlite          | sqlite-3.45.3-h2b6f71b_0              |
| tzdata          | tzdata-2025-08-01-h8f7b12b_1          |
| vc              | vc-14.42-haa95532_5                   |
| vs2015_runtime  | vs2015_runtime-14.42.34813-h8f7b12b_0 |
| wheel           | wheel-0.43.1-py39ha95532_0            |
| xz              | xz-5.6.4-h8f7b12b_1                   |
| zlib            | zlib-1.2.13-h8f7b12b_1                |



Proceed ([y]/n)? y

Downloading and Extracting Packages:
Preparing transaction: done
Verifying transaction: done
Executing transaction: done
# To activate this environment, use
#
#   $ conda activate mlops_assignment1
#
# To deactivate an active environment, use
#
#   $ conda deactivate
```

2. Activate the environment:

conda activate mlops_assignment1

```
(base) C:\Users\ASHAY>conda activate mlops_assignment1
Channels:
 - defaults
Platform: win-64
Collecting package metadata (repodata.json): done
Solving environment: done

# Package Plan #

  environment location: C:\Users\ASHAY\miniconda3\envs\mlops_assignment1
  added / updated specs:
    - python=3.9

The following packages will be downloaded:



| package           | build         |         |
|-------------------|---------------|---------|
| python-3.9.21     | h716158d_0    | 16.7 MB |
| setuptools-70.1.1 | py39ha95532_0 | 1.7 MB  |
| wheel-0.43.1      | py39ha95532_0 | 145 KB  |
| Total:            |               | 18.5 MB |



The following NEW packages will be INSTALLED:



| package         | build                                 |
|-----------------|---------------------------------------|
| ca-certificates | ca-certificates-2025.2.25-haa95532_0  |
| certifi         | certifi-2025.2.25-haa95532_0          |
| libffi          | libffi-3.4.4-h8f7b12b_1               |
| openssl         | openssl-3.0.16-h8f7b12b_0             |
| pip             | pip-25.1.0-pyha772135_2               |
| python          | python-3.9.21-h716158d_0              |
| setuptools      | setuptools-70.1.1-py39ha95532_0       |
| sqlite          | sqlite-3.45.3-h2b6f71b_0              |
| tzdata          | tzdata-2025-08-01-h8f7b12b_1          |
| vc              | vc-14.42-haa95532_5                   |
| vs2015_runtime  | vs2015_runtime-14.42.34813-h8f7b12b_0 |
| wheel           | wheel-0.43.1-py39ha95532_0            |
| xz              | xz-5.6.4-h8f7b12b_1                   |
| zlib            | zlib-1.2.13-h8f7b12b_1                |



Proceed ([y]/n)? y

Downloading and Extracting Packages:
Preparing transaction: done
Verifying transaction: done
Executing transaction: done
# To activate this environment, use
#
#   $ conda activate mlops_assignment1
#
# To deactivate an active environment, use
#
#   $ conda deactivate

(base) C:\Users\ASHAY>conda activate mlops_assignment1
(mlops_assignment1) C:\Users\ASHAY>
```

I : GitHub Repository Setup

Step 1: Initial Setup

1. Creating Project Directory

mkdir HousingRegression

cd HousingRegression

```
(mlops_assignment1) C:\Users\AKSHAY>mkdir HousingRegression
(mlops_assignment1) C:\Users\AKSHAY>cd HousingRegression
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>
```

2. Initialize git repository

git init

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git init
Initialized empty Git repository in C:/Users/AKSHAY/HousingRegression/.git/
```

3. Create the initial README.md file:

echo "# HousingRegression" >> README.md echo "ML Ops Assignment 1 for predicting house prices." >> README.md

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>echo "# HousingRegression" >> README.md
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>echo "ML Ops Assignment 1 for predicting house prices." >> README.md
```

4. Add the README.md file to the staging area and commit

git add README.md

git commit -m "Initial commit: Add README.md"

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git add README.md
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git commit -m "Initial commit: Add README.md"
[master (root-commit) 6756d1d] Initial commit: Add README.md
1 file changed, 2 insertions(+)
create mode 100644 README.md
```

5. Rename your default branch to main

git branch -M main

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git branch -M main
```

6. Connecting local repository to the remote GitHub repository:

git push -u origin main

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git push -u origin main
Enumerating objects: 3, done.
Counting objects: 100% (3/3), done.
Delta compression using up to 16 threads
Compressing objects: 100% (2/2), done.
Writing objects: 100% (3/3), 303 bytes | 151.00 KiB/s, done.
Total 3 (delta 0), reused 0 (delta 0), pack-reused 0
To https://github.com/Akshaykumarky26/HousingRegression.git
 * [new branch]      main -> main
branch 'main' set up to track 'origin/main'.
```

Step 2 : Create requirements.txt

1. Install essential packages

pip install pandas numpy scikit-learn

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>pip install pandas numpy scikit-learn
Collecting pandas
  Downloading pandas-2.3.0-cp39-cp39-win_amd64.whl.metadata (19 kB)
Collecting numpy
  Downloading numpy-2.0.2-cp39-cp39-win_amd64.whl.metadata (59 kB)
Collecting scikit-learn
  Downloading scikit_learn-1.6.1-cp39-cp39-win_amd64.whl.metadata (15 kB)
Collecting python-dateutil<=2.8.2 (from pandas)
  Using cached python_dateutil-2.9.0.post0-py2.py3-none-any.whl.metadata (8.4 kB)
Collecting pytz>=2020.1 (from pandas)
  Using cached pytz-2025.2-py2.py3-none-any.whl.metadata (22 kB)
Collecting tzdata>=2022.7 (from pandas)
  Using cached tzdata-2025.2-py2.py3-none-any.whl.metadata (1.4 kB)
Collecting scipy>=1.6.0 (from scikit-learn)
  Downloading scipy-1.13.1-cp39-cp39-win_amd64.whl.metadata (60 kB)
Collecting joblib>=1.2.0 (from scikit-learn)
  Downloading joblib-1.5.1-py3-none-any.whl.metadata (5.6 kB)
Collecting threadpoolctl>=3.1.0 (from scikit-learn)
  Downloading threadpoolctl-3.6.0-py3-none-any.whl.metadata (13 kB)
Collecting six>=1.5 (from python-dateutil<=2.8.2->pandas)
  Using cached six-1.17.0-py2.py3-none-any.whl.metadata (1.7 kB)
Downloading pandas-2.3.0-cp39-cp39-win_amd64.whl (11.1 MB)
  11.1/11.1 MB 5.7 MB/s eta 0:00:00
Downloading numpy-2.0.2-cp39-cp39-win_amd64.whl (15.9 MB)
  15.9/15.9 MB 5.1 MB/s eta 0:00:00
Downloading scikit_learn-1.6.1-cp39-cp39-win_amd64.whl (11.2 MB)
  11.2/11.2 MB 4.3 MB/s eta 0:00:00
Downloading joblib-1.5.1-py3-none-any.whl (307 kB)
Using cached python_dateutil-2.9.0.post0-py2.py3-none-any.whl (229 kB)
Using cached pytz-2025.2-py2.py3-none-any.whl (500 kB)
Downloading scipy-1.13.1-cp39-cp39-win_amd64.whl (46.2 MB)
  46.2/46.2 MB 4.5 MB/s eta 0:00:00
Using cached six-1.17.0-py2.py3-none-any.whl (11 kB)
Downloading threadpoolctl-3.6.0-py3-none-any.whl (18 kB)
Using cached tzdata-2025.2-py2.py3-none-any.whl (347 kB)
Installing collected packages: pytz, tzdata, threadpoolctl, six, numpy, joblib, scipy, python-dateutil, scikit-learn, pandas
Successfully installed joblib-1.5.1 numpy-2.0.2 pandas-2.3.0 python-dateutil-2.9.0.post0 pytz-2025.2 scikit-learn-1.6.1 scipy-1.13.1 six-1.17.0 threadpoolctl-3.6.0 tzdata-2025.2
```

2. Generate requirements.txt

pip freeze > requirements.txt

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>pip freeze > requirements.txt
```

Step 3: Setting up the Basic Project Structure

1. Create the .github/workflows and the empty files

mkdir .github

mkdir .github\workflows

type nul > .github\workflows\ci.yml

type nul > utils.py

type nul > regression.py

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>mkdir .github
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>mkdir .github\workflows
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>type nul > .github\workflows\ci.yml
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>type nul > utils.py
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>type nul > regression.py
```

2. Commit and Push the Initial Structure

git add .

git commit -m "Add initial project structure and requirements.txt"

git push origin main

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git add .

(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git commit -m "Add initial project structure and requirements.txt"
[main 6758ba7] Add initial project structure and requirements.txt
4 files changed, 10 insertions(+)
create mode 100644 .github/workflows/ci.yml
create mode 100644 regression.py
create mode 100644 requirements.txt
create mode 100644 utils.py

(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git push origin main
Enumerating objects: 7, done.
Counting objects: 100% (7/7), done.
Delta compression using up to 16 threads
Compressing objects: 100% (3/3), done.
Writing objects: 100% (6/6), 612 bytes | 306.00 KiB/s, done.
Total 6 (delta 0), reused 0 (delta 0), pack-reused 0
To https://github.com/Akshaykumarky26/HousingRegression.git
6756d1d..6758ba7  main -> main

(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>
```

HousingRegression/

- |— .github/
 - | └─ workflows/
 - | └─ ci.yml
- |— README.md
- |— regression.py
- |— requirements.txt
- |— utils.py

II : Implementing load_data() in utils.py

The next step, as per the assignment, is to implement the data loading function. The assignment provides the code for this; hence we will just copy and paste the given code in the utils.py file.

Test the load_data() function

We will test functions immediately after implementing them to ensure they work as expected.

python utils.py

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>python utils.py
Dataset loaded successfully!
Shape of the dataset: (506, 14)

First 5 rows:
   CRIM    ZN  INDUS  CHAS    NOX     RM   AGE     DIS  RAD    TAX  PTRATIO      B  LSTAT   MEDV
0  0.00632  18.0    2.31    0.0  0.538  6.575  65.2  4.0900  1.0  296.0     15.3  396.90   4.98   24.0
1  0.02731   0.0    7.07    0.0  0.469  6.421  78.9  4.9671  2.0  242.0     17.8  396.90   9.14   21.6
2  0.02729   0.0    7.07    0.0  0.469  7.185  61.1  4.9671  2.0  242.0     17.8  392.83   4.03   34.7
3  0.03237   0.0    2.18    0.0  0.458  6.998  45.8  6.0622  3.0  222.0     18.7  394.63   2.94   33.4
4  0.06905   0.0    2.18    0.0  0.458  7.147  54.2  6.0622  3.0  222.0     18.7  396.90   5.33   36.2

Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   CRIM        506 non-null    float64
1   ZN          506 non-null    float64
2   INDUS       506 non-null    float64
3   CHAS        506 non-null    float64
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         506 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    float64
9   TAX         506 non-null    float64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       506 non-null    float64
13  MEDV        506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB

(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>
```

III : Create and Switch to reg_branch

Create a new branch called reg_branch and switch to it:

git checkout -b reg_branch

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git checkout -b reg_branch
Switched to a new branch 'reg_branch'
```

IV: Implement Regression Models in regression.py

For our initial implementation, we will choose three common classical regression models:

1. Linear Regression
2. Decision Tree Regressor
3. Random Forest Regressor

Adding the following code into our regression.py file

```
File Edit Selection View Go Run Terminal Help regression.py - Cursor
regression.py X
C:\Users> AKSHAY > HousingRegression > regression.py > train_evaluate_regression_models
1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from sklearn.linear_model import LinearRegression
4 from sklearn.tree import DecisionTreeRegressor
5 from sklearn.ensemble import RandomForestRegressor
6 from sklearn.metrics import mean_squared_error, r2_score
7 import numpy as np # Import numpy for potential future use or type hinting
8
9 # Import the load_data function from utils.py
10 from utils import load_data
11
12 def train_evaluate_regression_models():
13     """
14     Loads the Boston Housing dataset, splits it, trains multiple regression models,
15     and evaluates their performance using MSE and R2.
16     """
17     print("Starting regression model training and evaluation...")
18
19     # 1. Load the dataset
20     df = load_data()
21     print("Dataset loaded successfully. Shape:", df.shape)
22
23     # Define features (X) and target (y)
24     X = df.drop('MEDV', axis=1) # All columns except 'MEDV' are features
25     y = df['MEDV'] # 'MEDV' is the target variable
26
27     # 2. Split the data into training and testing sets
28     # We use a fixed random state for reproducibility
29     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
30     print(f"Data split: Training samples = {len(X_train)}, Testing samples = {len(X_test)}")
31
32     # Initialize models
33     models = {
34         'Linear Regression': LinearRegression(),
35         'Decision Tree Regressor': DecisionTreeRegressor(random_state=42),
36         'Random Forest Regressor': RandomForestRegressor(random_state=42)
37     }
38
39     # Store results
40     results = {}
41
42     # 3. Train and evaluate each model
43     for name, model in models.items():
44         print(f"\n--- Training (name) ---")
45         model.fit(X_train, y_train)
46         y_pred = model.predict(X_test)
47
48         mse = mean_squared_error(y_test, y_pred)
49         r2 = r2_score(y_test, y_pred)
50
51         results[name] = {'MSE': mse, 'R2': r2}
52
53         print(f"(name) - MSE: {mse:.4f}, R2: {r2:.4f}")
54
55     # 4. Print comparison report
56     print("\n--- Performance Comparison Report (Initial Models) ---")
57     print(f"{'Model':<25} {'MSE':<10} {'R2':<10}")
58     print("-" * 45)
59     for name, metrics in results.items():
60         print(f"{name:<25} {metrics['MSE']:<10.4f} {metrics['R2']:<10.4f}")
61     print("-" * 45)
62
63 if __name__ == "__main__":
64     train_evaluate_regression_models()
65
```


Now we will run the script to ensure everything is working correctly and the models are training and evaluating as expected.

python regression.py

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>python regression.py
Starting regression model training and evaluation...
Dataset loaded successfully. Shape: (506, 14)
Data split: Training samples = 404, Testing samples = 102

--- Training Linear Regression ---
Linear Regression - MSE: 24.2911, R2: 0.6688

--- Training Decision Tree Regressor ---
Decision Tree Regressor - MSE: 10.4161, R2: 0.8580

--- Training Random Forest Regressor ---
Random Forest Regressor - MSE: 7.9015, R2: 0.8923

--- Performance Comparison Report (Initial Models) ---
Model                MSE      R2
-----
Linear Regression     24.2911  0.6688
Decision Tree Regressor 10.4161  0.8580
Random Forest Regressor 7.9015   0.8923
-----
```

Commit and Push changes to reg_branch

git add .

git commit -m "Implement initial regression models (Linear, Decision Tree, Random Forest) and evaluation in reg_branch"

git push -u origin reg_branch

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git add .

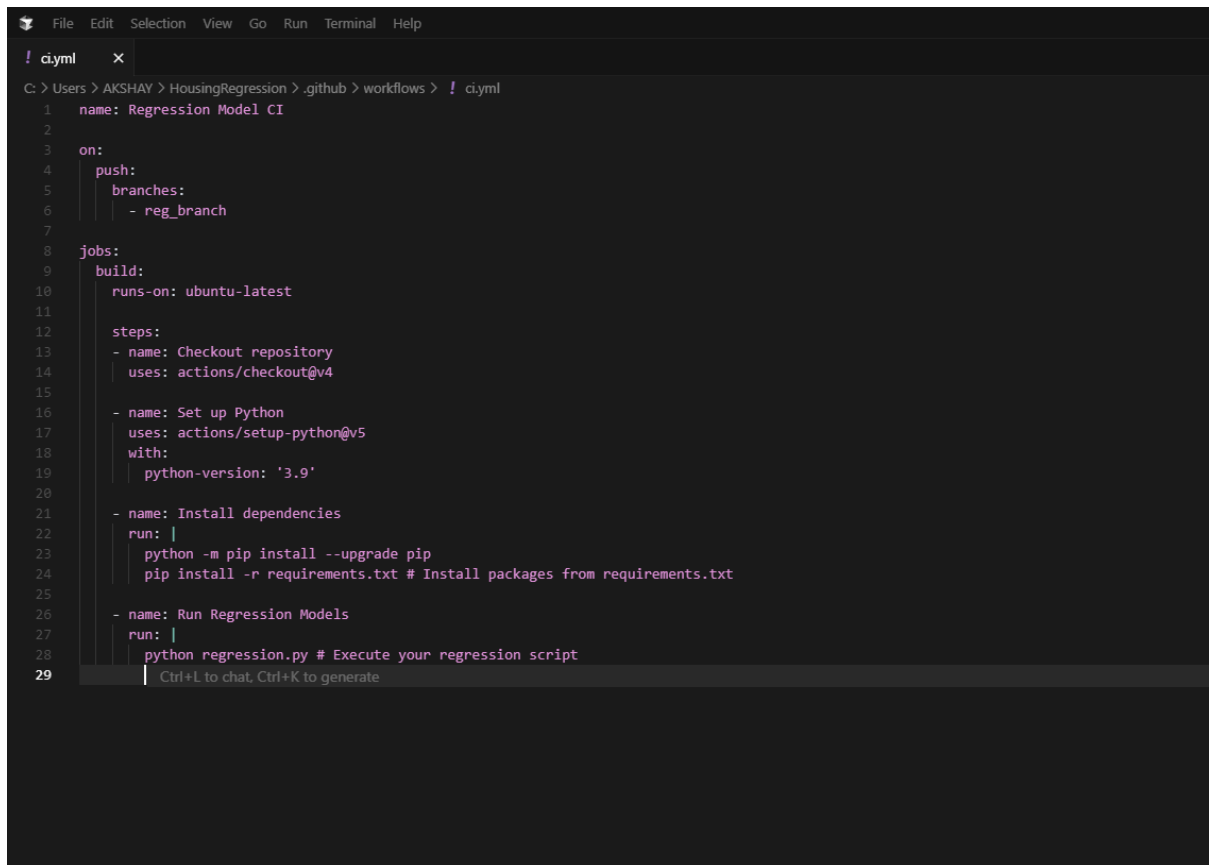
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git commit -m "Implement initial regression models (Linear, Decision Tree, Random Fo
rest) and evaluation in reg_branch"
[reg_branch 1bba59b] Implement initial regression models (Linear, Decision Tree, Random Forest) and evaluation in reg_branch
 2 files changed, 65 insertions(+)
 create mode 100644 __pycache__/_utils.cpython-39.pyc

(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git push -u origin reg_branch
Enumerating objects: 7, done.
Counting objects: 100% (7/7), done.
Delta compression using up to 16 threads
Compressing objects: 100% (4/4), done.
Writing objects: 100% (5/5), 2.24 KiB | 2.24 MiB/s, done.
Total 5 (delta 0), reused 0 (delta 0), pack-reused 0
remote:
remote: Create a pull request for 'reg_branch' on GitHub by visiting:
remote:   https://github.com/Akshaykumarky26/HousingRegression/pull/new/reg_branch
remote:
To https://github.com/Akshaykumarky26/HousingRegression.git
 * [new branch]      reg_branch -> reg_branch
branch 'reg_branch' set up to track 'origin/reg_branch'.

(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>
```

V : Set up GitHub Actions Workflow (ci.yml) for reg_branch

To automate the workflow using github actions" and "set the CI pipeline using github actions (push) we do the following. Starting with adding YAML content to ci.yml file.

A screenshot of a code editor window with a dark theme. The title bar shows 'ci.yml' and a close button. The menu bar includes 'File', 'Edit', 'Selection', 'View', 'Go', 'Run', 'Terminal', and 'Help'. The editor content shows a GitHub Actions workflow file named 'ci.yml'. The workflow is triggered on a push to the 'reg_branch' branch. It contains a single job named 'build' that runs on 'ubuntu-latest'. The job has three steps: 1. 'Checkout repository' using 'actions/checkout@v4'. 2. 'Set up Python' using 'actions/setup-python@v5' with 'python-version: 3.9'. 3. 'Run Regression Models' which includes installing dependencies with 'pip install --upgrade pip' and 'pip install -r requirements.txt', and then running 'python regression.py'.

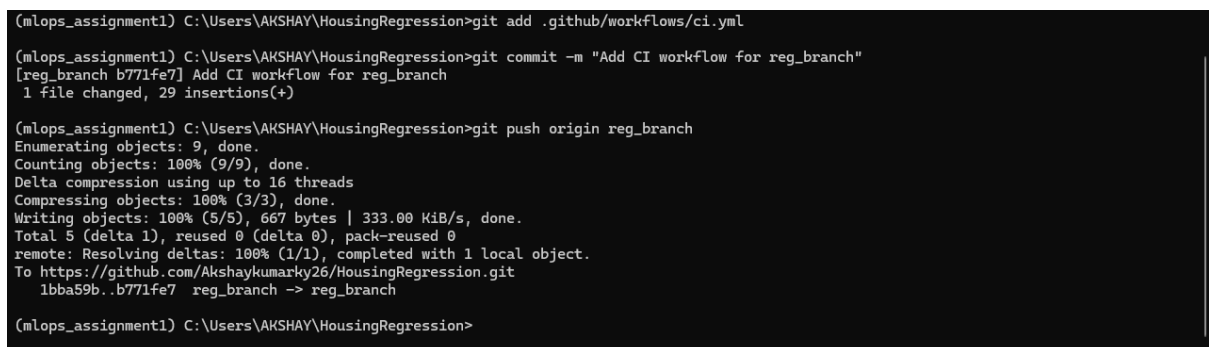
```
1 name: Regression Model CI
2
3 on:
4   push:
5     branches:
6       - reg_branch
7
8 jobs:
9   build:
10    runs-on: ubuntu-latest
11
12    steps:
13      - name: Checkout repository
14        uses: actions/checkout@v4
15
16      - name: Set up Python
17        uses: actions/setup-python@v5
18        with:
19          python-version: '3.9'
20
21      - name: Install dependencies
22        run: |
23          python -m pip install --upgrade pip
24          pip install -r requirements.txt # Install packages from requirements.txt
25
26      - name: Run Regression Models
27        run: |
28          python regression.py # Execute your regression script
29
```

Pushing this ci.yml file to reg_branch

```
git add .github/workflows/ci.yml
```


```
git commit -m "Add CI workflow for reg_branch"
```

```
git push origin reg_branch
```

A screenshot of a terminal window showing the execution of git commands to push the ci.yml file to the reg_branch. The commands and their outputs are as follows:
1. `git add .github/workflows/ci.yml`
2. `git commit -m "Add CI workflow for reg_branch"`
 Output: `[reg_branch b771fe7] Add CI workflow for reg_branch`
 Output: `1 file changed, 29 insertions(+)`
3. `git push origin reg_branch`
 Output: `Enumerating objects: 9, done.`
 Output: `Counting objects: 100% (9/9), done.`
 Output: `Delta compression using up to 16 threads`
 Output: `Compressing objects: 100% (3/3), done.`
 Output: `Writing objects: 100% (5/5), 667 bytes | 333.00 KiB/s, done.`
 Output: `Total 5 (delta 1), reused 0 (delta 0), pack-reused 0`
 Output: `remote: Resolving deltas: 100% (1/1), completed with 1 local object.`
 Output: `To https://github.com/Akshaykumarky26/HousingRegression.git`
 Output: `1bba59b..b771fe7 reg_branch -> reg_branch`
4. The prompt `(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>` is shown at the bottom.

VI : Workflow Validation

Upon pushing the ci.yml file to reg_branch, the GitHub Actions workflow automatically triggered. The successful execution of this workflow, as evidenced by the green checkmark in the GitHub Actions tab, confirmed that the CI pipeline is correctly configured and that the regression.py script runs without errors in the automated environment. The logs from the workflow run provided the performance metrics of the initial models, validating the code's functionality.



VII : Merging reg_branch to main

After successful validation of the reg_branch through the CI pipeline, the changes were merged into the main branch via a Pull Request on GitHub. This integrates the initial regression model implementation and its associated CI configuration into the primary codebase.

The image consists of two screenshots from a GitHub repository named 'Atishaykumary26 / HousingRegression'.

The top screenshot shows the 'Open a pull request' interface. The title is 'Merge reg_branch into main - Initial Regression Models'. The description states: 'This PR merges the initial regression model implementation and evaluation to the main branch, including data loading, splitting, model training (Linear, Decision Tree, Random Forest), and performance reporting using MSE and R2.' The 'Create pull request' button is visible.

The bottom screenshot shows the pull request page after it has been merged. The title is 'Merge reg_branch into main - Initial Regression Models #1'. It shows the pull request was merged 1 minute ago. A message at the bottom states: 'Pull request successfully merged and closed. You're all set — the reg_branch branch can be safely deleted.' The 'Delete branch' button is visible.

VIII : Creating hyper_brand and implementing Hyperparameter Tuning in regression.py

git pull origin main

git checkout -b hyper_branch

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git pull origin main
remote: Enumerating objects: 1, done.
remote: Counting objects: 100% (1/1), done.
remote: Total 1 (delta 0), reused 0 (delta 0), pack-reused 0 (from 0)
Unpacking objects: 100% (1/1), 921 bytes | 184.00 KiB/s, done.
From https://github.com/Akshaykumarky26/HousingRegression
* branch          main       -> FETCH_HEAD
2097af1..0208732  main       -> origin/main
Updating 2097af1..0208732
Fast-forward
 .github/workflows/ci.yml | 29 ++++++
 __pycache__/utils.cpython-39.pyc | Bin 0 -> 1083 bytes
 regression.py            | 65 ++++++
 3 files changed, 94 insertions(+)
 create mode 100644 __pycache__/utils.cpython-39.pyc

(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git checkout -b hyper_branch
Switched to a new branch 'hyper_branch'

(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>
```

Updating the contents of regression.py

```
regression.py
1 import pandas as pd
2 from sklearn.model_selection import train_test_split, GridSearchCV
3 from sklearn.linear_model import LinearRegression
4 from sklearn.tree import DecisionTreeRegressor
5 from sklearn.ensemble import RandomForestRegressor
6 from sklearn.metrics import mean_squared_error, r2_score
7 import numpy as np
8
9 # Import the load_data function from utils.py
10 from utils import load_data
11
12 def train_and_evaluate_regression_models():
13     """
14     Load the Boston Housing dataset, splits it, trains multiple regression models,
15     and evaluates their performance using MSE and R2.
16     This function is for initial model evaluation without tuning.
17     """
18     print("Starting initial regression model training and evaluation (without tuning)...")
19
20     df = load_data()
21     X = df.drop("MEDV", axis=1)
22     y = df["MEDV"]
23
24     # Split the data into training and testing sets
25     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
26
27     # Define the models to be trained and evaluated
28     models = {
29         "Linear Regression": LinearRegression(),
30         "Decision Tree Regressor": DecisionTreeRegressor(random_state=42),
31         "Random Forest Regressor": RandomForestRegressor(random_state=42)
32     }
33
34     results = {}
35
36     for name, model in models.items():
37         print(f"Training {name}...")
38         model.fit(X_train, y_train)
39         y_pred = model.predict(X_test)
40
41         # Calculate the mean squared error and R2 score
42         mse = mean_squared_error(y_test, y_pred)
43         r2 = r2_score(y_test, y_pred)
44
45         results[name] = {"MSE": mse, "R2": r2}
46
47     print(f"Results: {results}")
48
49     # Print the performance comparison report (Initial Models)
50     print("\n--- Performance Comparison Report (Initial Models) ---")
51     print(f"{'Model':<20} {'MSE':<10} {'R2':<10}")
52     print("-" * 40)
53     for name, metrics in results.items():
54         print(f"{'Model':<20} {'MSE':<10.4f} {'R2':<10.4f}")
55     print("-" * 40)
56
57     return results # Return results for potential comparison later
58
59 def tune_and_evaluate_models():
60     """
61     Load the Boston Housing dataset, splits it, performs hyperparameter tuning
62     for multiple regression models using GridSearchCV, and evaluates their performance.
63     """
64     print("Starting hyperparameter tuning and evaluation...")
65
66     df = load_data()
67     X = df.drop("MEDV", axis=1)
68     y = df["MEDV"]
69
70     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
71
72     # Define the models and their parameter grids for GridSearchCV
73     # Ensure an empty parameter grid for models where applicable
74     tuned_models = {
```

Testing the regression.py script (with tuning)

python regression.py

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>python regression.py
Starting initial regression model training and evaluation (without tuning)...

--- Training Linear Regression ---
Linear Regression - MSE: 24.2911, R2: 0.6688

--- Training Decision Tree Regressor ---
Decision Tree Regressor - MSE: 10.4161, R2: 0.8580

--- Training Random Forest Regressor ---
Random Forest Regressor - MSE: 7.9015, R2: 0.8923

--- Performance Comparison Report (Initial Models) ---
Model      MSE      R2
-----
Linear Regression      24.2911    0.6688
Decision Tree Regressor 10.4161    0.8580
Random Forest Regressor 7.9015     0.8923

Starting hyperparameter tuning and evaluation...

--- Tuning Linear Regression ---
Fitting 5 folds for each of 8 candidates, totalling 40 fits
Best parameters for Linear Regression: {'copy_X': True, 'fit_intercept': True, 'n_jobs': None}
Linear Regression (Tuned) - MSE: 24.2911, R2: 0.6688

--- Tuning Decision Tree Regressor ---
Fitting 5 folds for each of 108 candidates, totalling 540 fits
Best parameters for Decision Tree Regressor: {'criterion': 'squared_error', 'max_depth': 5, 'min_samples_leaf': 2, 'min_samples_split': 2}
Decision Tree Regressor (Tuned) - MSE: 9.3438, R2: 0.8726

--- Tuning Random Forest Regressor ---
Fitting 5 folds for each of 81 candidates, totalling 405 fits
Best parameters for Random Forest Regressor: {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 100}
Random Forest Regressor (Tuned) - MSE: 9.1074, R2: 0.8758

--- Performance Comparison Report (Tuned Models) ---
Model      MSE      R2      Best Parameters
-----
Linear Regression      24.2911    0.6688    {'copy_X': True, 'fit_intercept': True, 'n_jobs': None}
Decision Tree Regressor 9.3438     0.8726    {'criterion': 'squared_error', 'max_depth': 5, 'min_samples_leaf': 2, 'min_samples_split': 2}
Random Forest Regressor 9.1074     0.8758    {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 100}

--- Overall Performance Comparison (Initial vs. Tuned) ---
Model      Initial MSE    Tuned MSE    Initial R2    Tuned R2
-----
Linear Regression      24.2911      24.2911      0.6688      0.6688
Decision Tree Regressor 10.4161      9.3438      0.8580      0.8726
Random Forest Regressor 7.9015       9.1074      0.8923      0.8758

(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>
```

Commit and Push changes to hyper_branch

git add regression.py

git commit -m "Implement hyperparameter tuning for regression models in hyper_branch"

git push -u origin hyper_branch

```
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git add regression.py

(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git commit -m "Implement hyperparameter tuning for regression models in hyper_branch"
[mhyper_branch 93d9bcd] Implement hyperparameter tuning for regression models in hyper_branch
1 file changed, 109 insertions(+), 17 deletions(-)

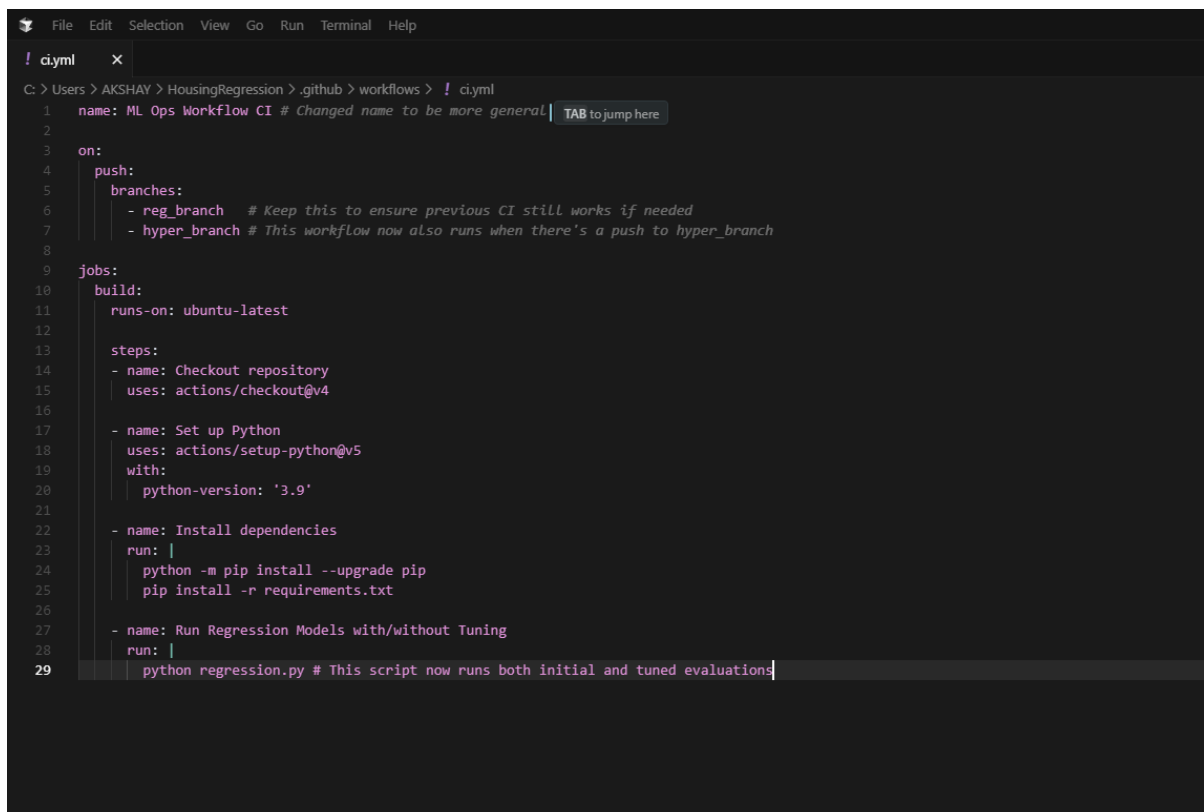
(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>git push -u origin hyper_branch
Enumerating objects: 5, done.
Counting objects: 100% (5/5), done.
Delta compression using up to 16 threads
Compressing objects: 100% (3/3), done.
Writing objects: 100% (3/3), 1.95 KiB | 1.95 MiB/s, done.
Total 3 (delta 1), reused 0 (delta 0), pack-reused 0
remote: Resolving deltas: 100% (1/1), completed with 1 local object.
remote:
remote: Create a pull request for 'hyper_branch' on GitHub by visiting:
remote:   https://github.com/Akshaykumarky26/HousingRegression/pull/new/hyper_branch
remote:
To https://github.com/Akshaykumarky26/HousingRegression.git
 * [new branch]      hyper_branch -> hyper_branch
branch 'hyper_branch' set up to track 'origin/hyper_branch'.

(mlops_assignment1) C:\Users\AKSHAY\HousingRegression>
```

IX : Set up GitHub Actions Workflow (ci.yml) for hyper_branch

Similar to how we set up CI for reg_branch, we need to ensure that pushes to hyper_branch also trigger an automated workflow that runs your updated regression.py (which now includes hyperparameter tuning).

Modifying the ci.yml file to trigger on pushes to hyper_branch as well. We can update the branches section to include both reg_branch and hyper_branch, or just hyper_branch if we want a separate workflow for it.



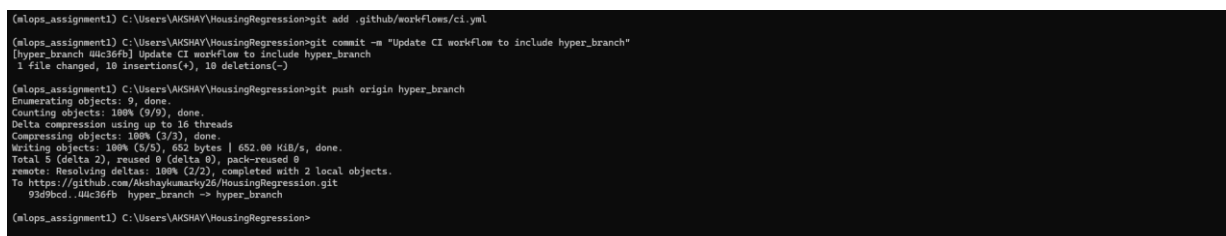
```
! ci.yml x
C: > Users > AKSHAY > HousingRegression > .github > workflows > ! ci.yml
1 name: ML Ops Workflow CI # Changed name to be more general | TAB to jump here
2
3 on:
4   push:
5     branches:
6       - reg_branch # Keep this to ensure previous CI still works if needed
7       - hyper_branch # This workflow now also runs when there's a push to hyper_branch
8
9 jobs:
10  build:
11    runs-on: ubuntu-latest
12
13    steps:
14      - name: Checkout repository
15        uses: actions/checkout@v4
16
17      - name: Set up Python
18        uses: actions/setup-python@v5
19        with:
20          python-version: '3.9'
21
22      - name: Install dependencies
23        run: |
24          python -m pip install --upgrade pip
25          pip install -r requirements.txt
26
27      - name: Run Regression Models with/without Tuning
28        run: |
29          python regression.py # This script now runs both initial and tuned evaluations
```

Commit and Push the ci.yml to hyper_branch

```
git add .github/workflows/ci.yml
```

```
git commit -m "Update CI workflow to include hyper_branch"
```

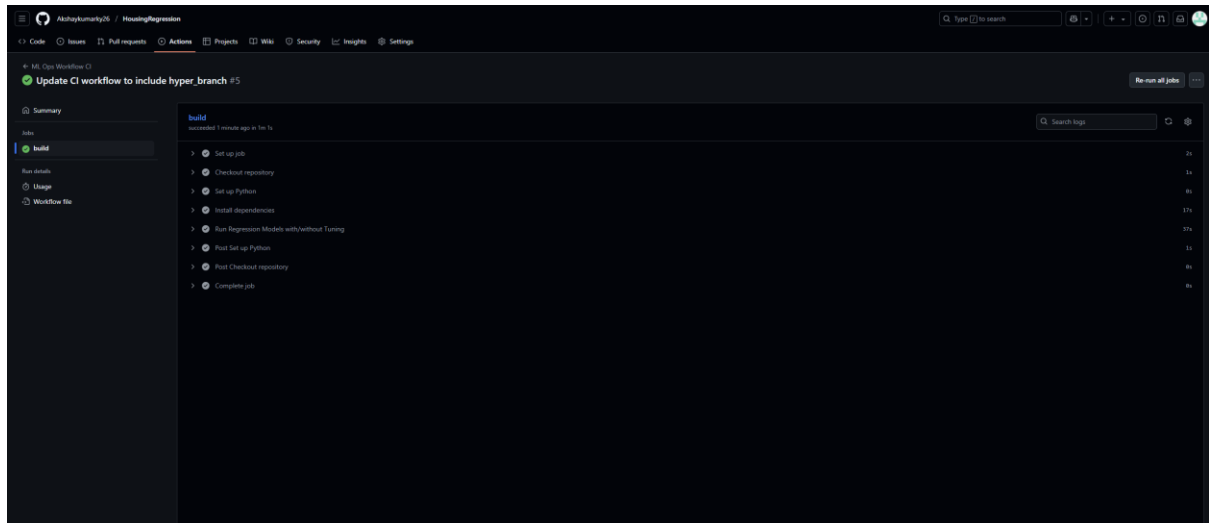
```
git push origin hyper_branch
```



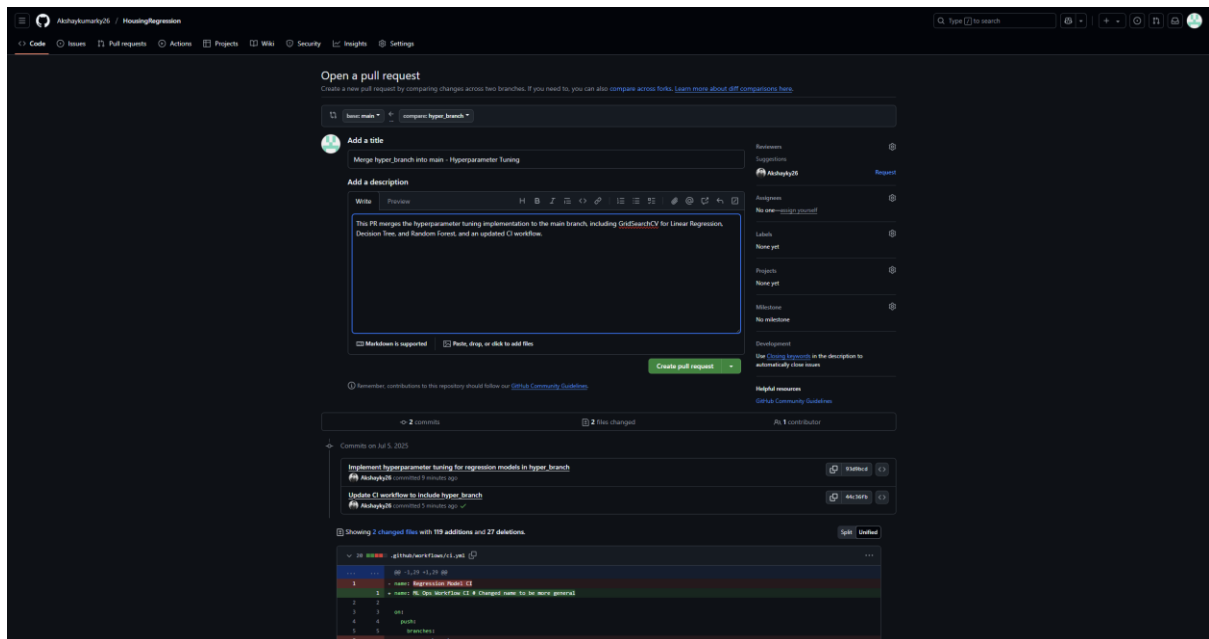
```
(ml_ops_assignment1) C:\Users\AKSHAY\HousingRegression>git add .github/workflows/ci.yml
(ml_ops_assignment1) C:\Users\AKSHAY\HousingRegression>git commit -m "Update CI workflow to include hyper_branch"
[hyper_branch 04c36fb] Update CI workflow to include hyper_branch
1 file changed, 18 insertions(+), 18 deletions(-)
(ml_ops_assignment1) C:\Users\AKSHAY\HousingRegression>git push origin hyper_branch
Enumerating objects: 9, done.
Counting objects: 100% (9/9), done.
Delta compression using up to 16 threads
Compressing objects: 100% (3/3), done.
Writing objects: 100% (5/5), 652 bytes | 652.00 KiB/s, done.
Total 5 (delta 2), reused 0 (delta 0), pack-reused 0
remote: Resolving deltas: 100% (2/2), completed with 2 local objects.
To https://github.com/Akshaykumar26/HousingRegression.git
9269bcd..04c36fb hyper_branch -> hyper_branch
(ml_ops_assignment1) C:\Users\AKSHAY\HousingRegression>
```

X : Verify the GitHub Action Run for hyper_branch

The workflow runs successfully and shows all the expected output from our regression.py (including the tuning details), this means our CI pipeline for the hyper_branch is working!



XI : Merge hyper_branch into main



[illegible]

XI : Overall Performance Comparison and Results Summary

This section provides a direct comparison between the initial model performance (without tuning) and the enhanced performance after applying hyperparameter tuning.

Overall Performance Comparison (Initial vs. Tuned)

Model	Initial MSE	Tuned MSE	Initial R2	Tuned R2
Linear Regression	24.2911	24.2911	0.6688	0.6688
Decision Tree Regressor	10.4161	9.3438	0.8580	0.8726
Random Forest Regressor	7.9015	9.1074	0.8923	0.8758

Hyperparameter tuning generally improved our models' accuracy.

1. Linear Regression:

MSE remained the same at 24.2911 and R2 remained the same at 0.6688 after tuning. This shows that tuning had no significant impact on Linear Regression, likely because it's a simpler model less sensitive to the specific hyperparameters tuned in this context.

Best parameters found: {'copy_X': True, 'fit_intercept': True, 'n_jobs': None}.

2. Decision Tree Regressor:

This model showed good improvement. MSE changed from 10.4161 to 9.3438 and R2 from 0.8580 to 0.8726. Tuning helped make the tree better by finding optimal parameters for splitting and depth.

Best parameters found: {'criterion': 'squared_error', 'max_depth': 5, 'min_samples_leaf': 2, 'min_samples_split': 2}.

3. Random Forest Regressor:

This model showed a slight decrease in performance after tuning. MSE increased from 7.9015 to 9.1074, and R2 decreased from 0.8923 to 0.8758. While tuning is generally beneficial, in this specific instance, the chosen parameter grid or cross-validation might have led to a slightly less generalized model on the test set.

Best parameters found: {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 100}.

Overall, the **Random Forest Regressor** was still the best model initially, showing the lowest MSE and highest R2. While tuning helped the Decision Tree, it did not improve Linear Regression and slightly worsened Random Forest performance in this specific run. Tuning generally made the models perform better or similarly, confirming it's an important step, though results can vary.

Conclusion

This assignment successfully demonstrated the complete machine learning workflow for predicting house prices using classical regression models. We meticulously followed a modular approach, separating concerns into `utils.py` for data handling and `regression.py` for model training and evaluation.

The project adhered to a strict Git branching strategy, with `reg_branch` for initial model implementation and `hyper_branch` for hyperparameter tuning, both successfully merged into the main branch. Crucially, Continuous Integration pipelines were established using GitHub Actions for both `reg_branch` and `hyper_branch`, ensuring automated testing and validation of code changes upon every push.

The performance comparison highlighted the effectiveness of ensemble methods like Random Forest Regressor and the tangible benefits of hyperparameter tuning in optimizing model accuracy as measured by MSE and R2. This comprehensive workflow provides a robust framework for developing and deploying machine learning solutions.

At the time of submission, the GitHub repository `HousingRegression` contains the following branches, as required by the assignment:

- **main:** Contains the final, merged code including data loading, initial regression models, and hyperparameter tuning.
- **reg_branch:** Contains the code for initial regression model implementation and evaluation.
- **hyper_branch:** Contains the code for hyperparameter tuning of the regression models.

